CLJDOC





Tap for Articles & Namespaces

Friendly SQL Functions

In Getting Started, we used execute! and execute-one! for all our SQL operations, except when we were reducing a result set. These functions (and plan) all expect a "connectable" and a vector containing a SQL string followed by any parameter values required.

A "connectable" can be a javax.sql.DataSource, a java.sql.Connection, or something that can produce a datasource (when get-datasource is called on it). It can also be a java.sql.PreparedStatement but we'll cover that a bit later...

Because string-building isn't always much fun, next.jdbc.sql also provides some "friendly" functions for basic CRUD operations:

- insert! and insert-multi! -- for inserting one or more rows into a table -- "Create",
- query -- an alias for execute! when using a vector of SQL and parameters -- "Read",
- update! -- for updating one or more rows in a table -- "Update",
- delete! -- for deleting one or more rows in a table -- "Delete".

as well as these more specific "read" operations:

- find-by-keys -- a query on one or more column values, specified as a hash map or WHERE clause,
- get-by-id -- a guery to return a single row, based on a single column value, usually the primary key.

These functions are described in more detail below. They are deliberately simple and intended to cover only

CLJDOC



- HoneySQL -- a composable DSL for creating SQL/parameter vectors from Clojure data structures
- seql -- a simplified EQL-inspired query language, built on next.jdbc (as of release 0.1.6)
- SQLingvo -- a composable DSL for creating SQL/parameter vectors
- Walkable -- full EQL query language support for creating SQL/parameter vectors

If you prefer to write your SQL separately from your code, take a look at HugSQL -- HugSQL documentation -- which has a next.jdbc adapter, as of version 0.5.1. See below for a "quick start" for using HugSQL with next.jdbc.

As of 1.3.925, aggregate-by-keys exists as a wrapper around find-by-keys that accepts the same options as find-by-keys and an aggregate SQL expression and it returns a single value (the aggregate). aggregate-by-keys accepts the same options as find-by-keys except that :columns may not be specified (since it is used to add the aggregate to the query).

insert!

Given a table name (as a keyword) and a hash map of column names and values, this performs a single row insertion into the database:

CLJDOC



If you have multiple rows (hash maps) to insert and they all have the same set of keys, you can use insert-multi! instead (see below), which will perform a single multi-row insertion, which will generally be faster.

insert-multi!

Given a table name (as a keyword), a vector of column names, and a vector of row value vectors, this performs a single multi-row insertion into the database:

All the row vectors must be the same length, and must match the number of columns specified.

Given a table name (as a keyword) and a vector of hash maps, this performs a single multi-row insertion into the database:

```
(sql/insert-multi! ds :address
```

CLJDOC



All the hash maps must have the same set of keys, so that the vector of hash maps can be converted to a vector of columns names and a vector of row value vectors, as above, so a single multi-row insertion can be performed.

If you wish to insert multiple hash maps that do not have identical keys, you need to iterate over insert! and insert one row at a time, which will generally be much slower.

Note: both of these expand to a single SQL statement with placeholders for every value being inserted --- for large sets of rows, this may exceed the limits on SQL string size and/or number of parameters for your JDBC driver or your database. Several databases have a limit of 1,000 parameter placeholders. Oracle does not support this form of multi-row insert, requiring a different syntax altogether.

Batch Insertion

As of release 1.2.790, you can specify :batch true in the options, which will use execute-batch! under the hood, instead of execute!, as follows:

```
(sql/insert-multi! ds :address
  [:name :email]
  [["Stella" "stella@artois.beer"]
   ["Waldo" "waldo@lagunitas.beer"]
   ["Aunt Sally" "sour@lagunitas.beer"]]
```

CLJDOC



```
["Stella" "stella@artois.beer"]
                      ["Waldo" "waldo@lagunitas.beer"]
                      ["Aunt Sally" "sour@lagunitas.beer"]]
                     {:return-keys true :return-generated-keys true})
;; and
(sql/insert-multi! ds :address
 [:name :email]
 [{:name "Stella", :email "stella@artois.beer"}
  {:name "Waldo", :email "waldo@lagunitas.beer"}
  {:name "Aunt Sally", :email "sour@lagunitas.beer"}]
 {:batch true})
;; equivalent to
(jdbc/execute-batch! ds
                     ["INSERT INTO address (name, email) VALUES (?,?)"
                      ["Stella" "stella@artois.beer"]
                      ["Waldo" "waldo@lagunitas.beer"]
                      ["Aunt Sally" "sour@lagunitas.beer"]]
                     {:return-keys true :return-generated-keys true})
```

Note: not all databases or drivers support returning generated keys like this -- see **Batched Parameters** for caveats and possible database-specific behaviors. You may need RETURNING * in your SQL instead.

query

Given a vector of SQL and parameters, execute it:

 $\textbf{com.github.seancorfield/next.jdbc} \quad 1.3.967$

CLJDOC



Note that the single argument form of <code>execute!</code> , taking just a <code>PreparedStatement</code> , is not supported by query .

update!

Given a table name (as a keyword), a hash map of columns names and values to set, and either a hash map of column names and values to match on or a vector containing a partial WHERE clause and parameters, perform an update operation on the database:

delete!

Given a table name (as a keyword) and either a hash map of column names and values to match on or a vector containing a partial WHERE clause and parameters, perform a delete operation on the database:

```
(sql/delete! ds :address {:id 8})
;; equivalent to
(sql/delete! ds :address ["id = ?" 8])
```

CLJDOC



Tap for Articles & Namespaces

find-by-keys

Given a table name (as a keyword) and either a hash map of column names and values to match on or a vector containing a partial WHERE clause and parameters, execute a query on the database:

While the hash map approach -- "query by example" -- is great for equality comparisons, sometimes you need other types of comparisons. For example, you might want to find all the rows where the email address ends in .beer:

```
(sql/find-by-keys ds :address ["email LIKE ?" "%.beer"])
;; equivalent to
(jdbc/execute! ds ["SELECT * FROM address WHERE email LIKE ?" "%.beer"])
```

Or you may want to find all the rows where the name is one of a specific set of values:

CLJDOC



The default behavior is to return all the columns in each row. You can specify a subset of columns to return using the columns option. It takes a vector and each element of the vector can be:

- a simple keyword representing the column name (:column-fn will be applied, if provided),
- a pair of keywords representing the column name and an alias (:column-fn will be applied to both, if provided),
- a pair consisting of a string and a keyword, representing a SQL expression and an alias (:column-fn will be applied to the alias, if provided).

Note: the SQL string provided for a column is copied exactly as-is into the generated SQL -- you are responsible for ensuring it is legal SQL!

find-by-keys supports an :order-by option which can specify a vector of column names to sort the results by. Elements may be column names or pairs of a column name and the direction to sort: :asc or :desc:

CLJDOC



find-by-keys also supports basic pagination with :offset and :fetch options which both accept numeric values and adds OFFSET ? ROWS FETCH NEXT ? ROWS ONLY to the generated query. To support MySQL and SQLite, you can specify :limit instead :fetch which adds LIMIT ? OFFSET ? to the generated query instead.

If you want to match all rows in a table -- perhaps with the pagination options in effect -- you can pass the keyword :all instead of either a hash map of column names and values or a vector containing a partial WHERE clause and parameters.

```
(sql/find-by-keys ds :address :all {:order-by [:id] :offset 5 :fetch 10})
;; equivalent to
(jdbc/execute! ds ["SELECT * FROM address ORDER BY id OFFSET ? ROWS FETCH NEXT ? ROWS ON
```

If no rows match, find-by-keys returns [], just like execute!.

aggregate-by-keys

Added in 1.3.925, this is a wrapper around find-by-keys that makes it easier to perform aggregate queries::

 $\textbf{com.github.seancorfield/next.jdbc} \quad 1.3.967$

CLJDOC



```
(get :next_jdbc_aggregate_123))
```

(where :next_jdbc_aggregate_123 is a unique alias generated by next.jdbc , derived from the aggregate expression string).

Note: the SQL string provided for the aggregate is copied exactly as-is into the generated SQL -- you are responsible for ensuring it is legal SQL!

get-by-id

Given a table name (as a keyword) and a primary key value, with an optional primary key column name, execute a guery on the database:

```
(sql/get-by-id ds :address 2)
;; equivalent to
(sql/get-by-id ds :address 2 {}) ; empty options map
;; equivalent to
(sql/get-by-id ds :address 2 :id {}) ; empty options map
;; equivalent to
(jdbc/execute-one! ds ["SELECT * FROM address WHERE id = ?" 2])
```

Note that in order to override the default primary key column name (of :id), you need to specify both the column name and an options hash map.

If no rows match, get-by-id returns nil, just like execute-one!.

CLJDOC



(unqualified) keywords provided. If you are trying to use a table name or column name that is a reserved name in SQL for your database, you will need to tell those functions to quote those names.

The namespace <code>next.jdbc.quoted</code> provides five functions that cover the most common types of entity quoting, and a modifier function for quoting dot-separated names (e.g., that include schemas):

- ansi -- wraps entity names in double quotes,
- mysql -- wraps entity names in back ticks,
- sql-server -- wraps entity names in square brackets,
- oracle -- an alias for ansi,
- postgres -- an alias for ansi.
- schema -- wraps a quoting function to support dbo.table style entity names.

These quoting functions can be provided to any of the friendly SQL functions above using the :table-fn and :column-fn options, in a hash map provided as the (optional) last argument in any call. If you want to provide your own entity naming function, you can do that:

```
(defn snake-case [s] (str/replace s #"-" "_"))
(sql/insert! ds :my-table {:some "data"} {:table-fn snake-case})
```

next.jdbc provides snake-kebab-opts and unqualified-snake-kebab-opts which are hash maps containing :column-fn and :table-fn that use the ->snake_case function from the camel-snake-kebab

CLJDOC




```
;; on any results, e.g., turning MySQL's :GENERATED_KEY into :generated-key (sql/insert! ds :my-table {:some "data"} jdbc/snake-kebab-opts)
```

Note: The entity naming function is passed a string, the result of calling name on the keyword passed in. Also note that the default quoting functions do not handle schema-qualified names, such as dbo.table_name -- sql-server would produce [dbo.table_name] from that. Use the schema function to wrap the quoting function if you need that behavior, e.g., {:table-fn (schema sql-server)} which would produce [dbo].[table_name].

HugSQL Quick Start

Here's how to get up and running quickly with <code>next.jdbc</code> and HugSQL. For more detail, consult the HugSQL documentation. Add the following dependencies to your project (in addition to <code>com.github.seancorfield/next.jdbc</code> and whichever JDBC drivers you need):

```
com.layerware/hugsql-core {:mvn/version "0.5.3"}
com.layerware/hugsql-adapter-next-jdbc {:mvn/version "0.5.3"}
```

Check the HugSQL documentation for the latest versions to use!

Write your SQL in .sql files that are on the classpath (somewhere under src or resources). For our purposes, assume a SQL file db/example.sql containing your first set of definitions. In your namespace, add these require s:

```
[hugsql.core :as hugsql]
```

CLJDOC



At program startup you'll need to call these functions (either at the top-level of your namespace on history your initialization function):

Those calls will add function definitions to that namespace based on what is in the sql files. Now set up your db-spec and datasource as usual with next.jdbc:

```
(def db-spec {:dbtype "h2:mem" :dbname "example"}) ; assumes H2 driver in deps.edn
(def ds (jdbc/get-datasource db-spec))
```

Borrowing from Princess Bride examples from the HugSQL documentation, you can now do things like this:

```
(create-characters-table ds)
;;=> [#:next.jdbc{:update-count 0}]
(insert-character ds {:name "Westley", :specialty "love"})
;;=> 1
```

CLJDOC



✓ Getting Started

Tips & Tricks >

Can you improve this documentation? These fine people already did: Sean Corfield, alexandrkozyrev & Valtteri Harmainen

Edit on GitHub