CLJDOC



Tap for Articles & Namespaces

Tips & Tricks

This page contains various tips and tricks that make it easier to use <code>next.jdbc</code> with a variety of databases. It is mostly organized by database, but there are a few that are cross-database and those are listed first.

CLOB & BLOB SQL Types

Columns declared with the CLOB or BLOB SQL types are typically rendered into Clojure result sets as database-specific custom types but they should implement java.sql.Clob or java.sql.Blob (as appropriate). In general, you can only read the data out of those Java objects during the current transaction, which effectively means that you need to do it either inside the reduction (for plan) or inside the result set builder (for execute! or execute-one!). If you always treat these types the same way for all columns across the whole of your application, you could simply extend next.jdbc.result-set/ReadableColumn to java.sql.Clob (and/or java.sql.Blob). Here's an example for reading CLOB into a String:

```
(extend-protocol rs/ReadableColumn
 java.sql.Clob
 (read-column-by-label [^java.sql.Clob v _]
   (with-open [rdr (.getCharacterStream v)] (slurp rdr)))
 (read-column-by-index [^java.sql.Clob v _2 _3]
   (with-open [rdr (.getCharacterStream v)] (slurp rdr))))
```

There is a helper in next.jdbc.result-set to make this easier -- clob->string:

CLJDOC



```
(rs/clob->string v))
(read-column-by-index [^java.sql.Clob v _2 _3]
  (rs/clob->string v)))
```

As noted in Result Set Builders, there is also clob-column-reader that can be used with the as-*-adapter result set builder functions.

No helper or column reader is provided for BLOB data since it is expected that the semantics of any given binary data will be application specific. For a raw byte[] you could probably use:

```
(.getBytes v 1 (.length v)) ; BLOB has 1-based byte index!
```

Consult the java.sql.Blob documentation for more ways to process it.

Note: the standard MySQL JDBC driver seems to return BLOB data as byte[] instead of java.sql.Blob.

Exceptions

A lot of JDBC operations can fail with an exception. JDBC 4.0 has a well-defined hierarchy of exception types and you can often catch a specific type of exception to do useful handling of various error conditions that you might "expect" when working with a database.

A good example is SQLIntegrityConstraintViolationException which typically represents an index/key constraint violation such as a duplicate primary key insertion attempt.

However, like some other areas when dealing with JDBC, the reality can be very database-specific. Some

CLJDOC



The java.sql.SQLException class provides .getErrorCode() and .getSQLState() methods but the values returned by those are explicitly vendor-specific (error code) or only partly standardized (state). In theory, the SQL state should follow either the X/Open (Open Group) or ANSI SQL 2003 conventions, both of which were behind paywalls(!). The most complete public listing is probably the IBM DB2 SQL State document. See also this Stack Overflow post about SQL State for more references and links. Not all database drivers follow either of these conventions for SQL State so you may still have to consult your vendor's specific documentation.

All of this makes writing *generic* error handling, that works across multiple databases, very hard indeed. You can't rely on the JDBC SQLException hierarchy; you can sometimes rely on a subset of SQL State values.

Handling Timeouts

JDBC provides a number of ways in which you can decide how long an operation should run before it times out. Some of these timeouts are specified in seconds and some are in milliseconds. Some are handled via connection properties (or JDBC URL parameters), some are handled via methods on various JDBC objects.

Here's how to specify various timeouts using next.jdbc:

- connectTimeout -- can be specified via the "db-spec" hash map or in a JDBC URL, it is the number of **milliseconds** that JDBC should wait for the initial (socket) connection to complete. Database-specific (may be MySQL only?).
- loginTimeout -- can be set via .setLoginTimeout() on a DriverManager or DataSource, it is the number of **seconds** that JDBC should wait for a connection to the database to be made. next.jdbc exposes this on the javax.sql.DataSource object it reifies from calling get-datasource on a "db-spec" hash map or JDBC URL string.

CLJDOC



passed to any function that may construct a statement of Preparedstatement.

• socketTimeout -- can be specified via the "db-spec" hash map or in a JDBC URL, it is the number of milliseconds that JDBC should wait for socket operations to complete. Database-specific (MS SQL Server and MySQL support this, other databases may too).

Examples:

```
;; connectTimeout / socketTimeout via db-spec:
(def db-spec {:dbtype "mysql" :dbname "example" :user "root" :password "secret"
              ;; milliseconds:
              :connectTimeout 60000 :socketTimeout 30000}))
;; socketTimeout via JDBC URL:
(def db-url (str "jdbc:sqlserver://localhost;user=sa;password=secret"
                 ;; milliseconds:
                 ";database=model;socketTimeout=10000"))
;; loginTimeout via DataSource:
(def ds (jdbc/get-datasource db-spec))
(.setLoginTimeout ds 20); seconds
;; queryTimeout via options:
(jdbc/execute! ds ["select * from some_table"] {:timeout 5}) ; seconds
;; queryTimeout via method call:
(let [ps (jdbc/prepare ds ["select * from some_table"])]
 (.setQueryTimeout ps 10); seconds
 (jdbc/execute! ps))
```

CLJDOC



overhead of realizing rows from the Resultset into Clojure data structures if your reducing function uses only functions that get column values by name. If you perform any function on the row that would require an actual hash map or a sequence, the row will be realized into a full Clojure hash map via the builder function passed in the options (or via next.jdbc.result-set/as-maps by default).

One of the benefits of reducing over plan is that you can stream very large result sets, very efficiently, without having the entire result set in memory (assuming your reducing function doesn't build a data structure that is too large!). See the tips below on **Streaming Result Sets**. If you want to process a plan result purely for side-effects, without producing a result, you can use run! instead of reduce:

```
(run! process-row (jdbc/plan ds ...))
```

run! is based on reduce and process-row here takes just one argument -- the row -- rather than the usual reducing function that takes two.

The result of plan is also foldable in the clojure.core.reducers sense. While you could use execute! to produce a vector of fully-realized rows as hash maps and then fold that vector (Clojure's vectors support fork-join parallel reduce-combine), that wouldn't be possible for very large result sets. If you fold the result of plan, the result set will be partitioned and processed using fork-join parallel reduce-combine. Unlike reducing over plan, each row is realized into a Clojure data structure and each batch is forked for reduction as soon as that many rows have been realized. By default, fold 's batch size is 512 but you can specify a different value in the 4-arity call. Once the entire result set has been read, the last (partial) batch is forked for reduction. The combining operations are forked and interleaved with the reducing operations, so the order (of forked tasks) is batch-1, batch-2, combine-1-2, batch-3, combine-1&2-3, batch-4, combine-1&2&3-4, etc. The amount of parallelization you get will depend on many factors including the number of processors, the speed of your reducing function, the speed of your combining function, and the

CLJDOC



Tap for Articles & Namespaces

Times, Dates, and Timezones

Working with dates and timezones in databases can be confusing, as you are working at the intersection between the database, the JDBC library and the date library that you happen to be using. A good rule of thumb is to keep timezone-related logic as simple as possible. For example, with Postgres we recommend always storing dates in a Postgres TIMESTAMP (without time zone) column, storing all such timestamps in UTC, and applying your time zone logic separately using application logic. The TIMESTAMP WITH TIME ZONE column type in Postgres stores its date in UTC anyhow, and applications that need to deal with time zones typically require richer functionality than simply adjusting the time zone to wherever the database happens to be hosted. Treat time zone related logic as an application concern, and keep stored dates in UTC.

For example, for a developer using clojure.java-time, saving (java-time/instant) in a timestamp column (and doing any timezone adjustment elsewhere) is a good way to minimize long term confusion.

Original text contributed by Denis McCarthy; in addition: I generally recommend not only using UTC everywhere but also setting your database and your servers to all be in the UTC timezones, to avoid the possibly of incorrect date/time translations -- Sean Corfield.

MS SQL Server

In MS SQL Server, the generated key from an insert comes back as : GENERATED_KEYS .

By default, you won't get table names as qualifiers with Microsoft's JDBC driver (you might with the jTDS drive -- I haven't tried that recently). See this MSDN forum post about .getTableName() for details. According to one of the answers posted there, if you specify :result-type and :concurrency in the

CLJDOC



MS SQL Server supports execution of multiple statements when surrounded by begin / end and can return multiple result sets, when requested via :multi-rs true on execute!.

```
(jdbc/execute! db-spec ["begin select * from table1; select * from table2; end"] {:multi
;; vector of result sets:
=> [[{.. table1 row ..} {.. table1 row ..}]
     [{.. table2 row ..} {.. table2 row ..} {..}]]
```

Batch Statements

Even when using <code>next.jdbc/execute-batch!</code>, Microsoft's JDBC driver will still send multiple insert statements to the database unless you specify <code>:useBulkCopyForBatchInsert true</code> as part of the db-spec hash map or JDBC URL when the datasource is created.

To use this feature your Microsoft's JDBC driver should be at least version 9.2 and you can use only limited set of data types. For example if you use inst to bulk insert smalldatetime value driver will revert to old (slow) behavior. For more details see Using bulk copy API for batch insert operation and Release notes for JDBC drivers.

MySQL & MariaDB

In MySQL, the generated key from an insert comes back as :GENERATED_KEY . In MariaDB, the generated key from an insert comes back as :insert_id .

MySQL generally stores tables as files so they are case-sensitive if your O/S is (Linux) or case-insensitive if your O/S is not (Mac, Windows) but the column names are generally case-insensitive. This can matter when

CLJDOC



טבטמעטב טו נוווס.

It's also worth noting that column comparisons are case-insensitive so where foo = 'BAR' will match "bar" or "BAR" etc.

MySQL has a connection option, :allowMultiQueries true, that allows you to pass multiple SQL statements in a single operation and can return multiple result sets, when requested via :multi-rs true.

```
(def db-spec {:dbtype "mysql" .. :allowMultiQueries true})
;; equivalent to allowMultiQueries=true in the JDBC URL
(jdbc/execute! db-spec ["select * from table1; select * from table2"] {:multi-rs true})
;; vector of result sets:
=> [[{.. table1 row ..} {.. table1 row ..}]
     [{.. table2 row ..} {.. table2 row ..} {..}]]
```

Compare this with MS SQL Server above: MySQL does not support begin / end here. This is not the default behavior because allowing multiple statements in a single operation is generally considered a bit of a risk as it can make it easier for SQL injection attacks to be performed.

Batch Statements

Even when using next.jdbc/execute-batch!, MySQL will still send multiple statements to the database unless you specify :rewriteBatchedStatements true as part of the db-spec hash map or JDBC URL when the datasource is created.

Streaming Result Sets

CLJDOC



Tap for Articles & Namespaces

Note: it's possible that other options may be required as well -- I have not verified this yet -- see, for example, the additional options PostgreSQL requires, below.

Oracle

Ah, dear old Oracle! Over the years of maintaining clojure.java.jdbc and now next.jdbc, I've had all sorts of bizarre and non-standard behavior reported from Oracle users. The main issue I'm aware of with next.jdbc is that Oracle's JDBC drivers all return an empty string from ResultSetMetaData.getTableName() so you won't get qualified keywords in the result set hash maps. Sorry!

An important performance issue to be aware of with Oracle's JDBC driver is that the default fetch size is just 10 records. If you are working with large datasets, you will either need to either specify :prefetch in your db-spec hash map with a suitable value (say 1,000 or larger), or specify &prefetch= in your JDBC URL string. If you want to keep the default, you can change it on a per-statement basis by specifying :fetchsize as an option to execute! etc.

If you are using the 10g or later JDBC driver and you try to execute DDL statements that include SQL entities that start with a : (such as :new or :old), they will be treated as bindable parameter references if you use a PreparedStatement to execute them. Since that's the default for execute! etc, it means that you will likely get an error like the following:

Missing IN or OUT parameter at index:: 1

You will need to use next.jdbc.prepare/statement to create a Statement object and then call

CLJDOC



Tap for Articles & Namespaces

PUSIGIEDUL

As you can see in this section (and elsewhere in this documentation), the PostgreSQL JDBC driver has a number of interesting quirks and behaviors that you need to be aware of. Although accessing PostgreSQL via JDBC is the most common approach, there is also a non-JDBC Clojure/Java driver for PostgreSQL called PG2 which supports JSON operations natively (see below for what's required for JDBC), as well as supporting Java Time natively (see the section above about Times, Dates, and Timezones), and it also quite a bit faster than using JDBC.

When you use :return-keys true with execute! or execute-one! (or you use insert!), PostgreSQL returns the entire inserted row (unlike nearly every other database that just returns any generated keys!).

The default result set builder for next.jdbc is as-qualified-maps which uses the .getTableName() method on ResultSetMetaData to qualify the columns in the result set. While some database drivers have this information on hand from the original SQL operation, PostgreSQL's JDBC driver does not and it performs an extra SQL query to fetch table names the first time this method is called for each query. If you want to avoid those extra queries, and you can live with unqualified column names, you can use asunqualified-maps as the result set builder instead.

If you have a query where you want to select where a column is IN a sequence of values, you can use col = ANY(?) with a native array of the values instead of IN (?,?,?,,,?) and a sequence of values. **Be** aware of PostgreSQL bug 17822 which can cause pathological performance when the array has a single element! If you think you might have a single-element array, consider using UNNEST and IN instead.

What does this mean for your use of next.jdbc? In plan, execute!, and execute-one!, you can use col = ANY(?) in the SQL string and a single primitive array parameter, such as (int-array [1 2 3 4]).

CLJDOC



Tap for Articles & Namespaces

Batch Statements

Even when using next.jdbc/execute-batch!, PostgreSQL will still send multiple statements to the database unless you specify :reWriteBatchedInserts true as part of the db-spec hash map or JDBC URL when the datasource is created.

Streaming Result Sets

You can get PostgreSQL to stream very large result sets (when you are reducing over plan) by setting the following options:

- :auto-commit false -- when opening the connection
- :fetch-size 4000, :concurrency :read-only, :cursors :close, :result-type :forward-only -- when running plan (or when creating a PreparedStatement).

Working with Arrays

ResultSet protocol extension to read SQL arrays as Clojure vectors.

CLJDOC




```
(extend-protocol rs/keadablecolumn
Array

(read-column-by-label [^Array v _] (vec (.getArray v)))

(read-column-by-index [^Array v _ _] (vec (.getArray v))))
```

Insert and read vector example:

```
create table example(
  tags varchar[]
);

(execute-one! db-spec
  ["insert into example(tags) values (?)"
      (into-array String ["tag1" "tag2"])])
```

Note: PostgreSQL JDBC driver supports only 7 primitive array types, but not array types like UUID[] - PostgreSQL™ Extensions to the JDBC API.

Working with Date and Time

(execute-one! db-spec

["select * from example limit 1"])

;; => #:example{:tags ["tag1" "tag2"]}

CLJDOC



In addition, if you want java.time.Instant, java.time.LocalDate, and java.time.LocalDateTime to be automatically converted to SQL data types, requiring next.jdbc.date-time will enable those as well (by extending SettableParameter for you).

next.jdbc.date-time also includes functions that you can call at application startup to extend ReadableColumn to either return java.time.Instant or java.time.LocalDate/java.time.LocalDateTime (as well as a function to restore the default behavior of returning java.sql.Date and java.sql.Timestamp).

Working with Interval

Postgres has a nonstandard SQL type Interval that is implemented in the Postgres driver as the org.postgresql.util.PGInterval type. In many cases you would want to work with intervals as java.time.Duration type by default.

You can support Duration instances by extending SettableParameter to the java.time.Duration type. Conversely you can support converting PGIntervals back to Durations by extending ReadableColumn to the org.postgresql.util.PGInterval type.

```
(import '[org.postgresql.util PGInterval])
(import '[java.sql PreparedStatement])
(import '[java.time Duration])
(require '[next.jdbc.result-set :as rs])
(require '[next.jdbc.prepare :as p])

(defn ->pg-interval
    "Takes a Dudration instance and converts it into a PGInterval
```

CLJDOC



```
(extend-protocol p/SettableParameter
 ;; Convert durations to PGIntervals before inserting into db
 java.time.Duration
 (set-parameter [^java.time.Duration v ^PreparedStatement s ^long i]
   (.setObject s i (->pg-interval v))))
(defn <-pg-interval</pre>
  "Takes a PGInterval instance and converts it into a Duration
  instance. Ignore sub-second units."
 [^org.postgresql.util.PGInterval interval]
 (-> Duration/ZERO
      (.plusSeconds (.getSeconds interval))
      (.plusMinutes (.getMinutes interval))
      (.plusHours (.getHours interval))
      (.plusDays (.getDays interval))))
(extend-protocol rs/ReadableColumn
 ;; Convert PGIntervals back to durations
 org.postgresgl.util.PGInterval
 (read-column-by-label [^org.postgresgl.util.PGInterval v _]
   (<-pq-interval v))</pre>
 (read-column-by-index [^org.postgresql.util.PGInterval v _2 _3]
   (<-pg-interval v)))</pre>
```

Working with Enumerated Types

PostgreSQL has a SQL extension for defining enumerated types and the default set-parameter

CLJDOC



```
CREATE TYPE language AS ENUM('en','fr','de');

CREATE TABLE person (
    ...
    speaks language NOT NULL,
    ...
);

(require '[next.jdbc.sql :as sql]
        '[next.jdbc.types :refer [as-other]])

(sql/insert! ds :person {:speaks (as-other "fr")})
```

That call produces a vector ["fr"] with metadata that implements set-parameter such that .setObject() is called with java.sql.Types/OTHER which allows PostgreSQL to "convert" the string "fr" to the corresponding language enumerated type value.

Working with JSON and JSONB

PostgreSQL has good support for storing, querying and manipulating JSON data. Basic Clojure data structures (lists, vectors, and maps) transform pretty well to JSON data. With a little help next.jdbc can automatically convert Clojure data to JSON and back for us.

Note: some PostgreSQL JSONB operators have a ? in them which conflicts with the standard parameter placeholder in SQL. You can write the JSONB operators by doubling up the ?, e.g., ??! instead of just ?!. See PostgreSQL JSONB operators for more detail.

CLJDOC




```
(require '[]sonista.core :as ]son])

;; :decode-key-fn here specifies that JSON-keys will become keywords:
(def mapper (json/object-mapper {:decode-key-fn keyword}))
(def ->json json/write-value-as-string)
(def <-json #(json/read-value % mapper))</pre>
```

Next we create helper functions to transform Clojure data to and from PostgreSQL Objects containing JSON:

```
(import '(org.postgresql.util PGobject))
(defn ->pgobject
 "Transforms Clojure data to a PGobject that contains the data as
 JSON. PGObject type defaults to `jsonb` but can be changed via
 metadata key `:pqtype`"
 [x]
 (let [pgtype (or (:pgtype (meta x)) "jsonb")]
   (doto (PGobject.)
      (.setType pgtype)
      (.setValue (->json x))))
(defn <-pgobject</pre>
 "Transform PGobject containing `json` or `jsonb` value to Clojure data."
 [^PGobject v]
 (let [type (.getType v)
       value (.getValue v)]
   (if (#{"jsonb" "json"} type)
      (some-> value <-json (with-meta {:pgtype type}))</pre>
```

CLJDOC



protocols to make the conversion between clojure data and PGobject JSON automatic:

```
(require '[next.jdbc.prepare :as prepare])
(require '[next.jdbc.result-set :as rs])
(import '[java.sql PreparedStatement])
(set! *warn-on-reflection* true)
;; if a SQL parameter is a Clojure hash map or vector, it'll be transformed
;; to a PGobject for JSON/JSONB:
(extend-protocol prepare/SettableParameter
 clojure.lang.IPersistentMap
 (set-parameter [m ^PreparedStatement s i]
   (.setObject s i (->pgobject m)))
 clojure.lang.IPersistentVector
 (set-parameter [v ^PreparedStatement s i]
   (.setObject s i (->pgobject v))))
;; if a row contains a PGobject then we'll convert them to Clojure data
;; while reading (if column is either "json" or "jsonb" type):
(extend-protocol rs/ReadableColumn
 org.postgresql.util.PGobject
 (read-column-by-label [^org.postgresql.util.PGobject v _]
   (<-pgobject v))</pre>
 (read-column-by-index [^org.postgresql.util.PGobject v _2 _3]
   (<-pgobject v)))</pre>
```

CLJDOC



We can now insert Clojure data into json and jsonb fields:

```
(require '[next.jdbc :as jdbc])
(require '[next.jdbc.sql :as sql])
(def db { ...db-spec here... })
(def ds (jdbc/get-datasource db))
(def test-map
 {:some-key "some val" :nested {:a 1} :null-val nil :vector [1 2 3]})
(def data1
 {:doc_jsonb test-map
   :doc_json (with-meta test-map {:pgtype "json"})})
(sql/insert! ds :demo data1)
(def test-vector
    [{:a 1} nil 2 "lalala" []])
(def data2
   {:doc_jsonb test-vector
```

CLJDOC



And those columns are nicely transformed into Clojure data when querying:

```
(sql/get-by-id ds :demo 1)
=> #:demo{:id 1,
          :doc_json
          {:some-key "some val",
           :nested {:a 1},
           :vector [1 2 3],
           :null-val nil},
          :doc_jsonb
          {:some-key "some val",
           :nested {:a 1},
           :vector [1 2 3],
           :null-val nil}}
(sql/get-by-id ds :demo 2)
=> #:demo{:id 2,
          :doc_json [{:a 1} nil 2 "lalala" []],
          :doc_jsonb [{:a 1} nil 2 "lalala" []]}
;; Query by value of JSON field 'some-key'
(sql/query ds [(str "select id, doc_jsonb::json->'nested' as foo"
                    " from demo where doc_jsonb::json->>'some-key' = ?")
               "some val"])
=> [{:demo/id 1, :foo {:a 1}}]
```

Using HoneySQL with JSON and JSONB

CLJDOC



in the noneyear accumentation for more actains.

JSON or JSONB?

- A json column stores JSON data as strings (reading and writing is fast but manipulation is slow, field order is preserved)
- A jsonb column stores JSON data in binary format (manipulation is significantly faster but reading and writing is a little slower)

If you're unsure whether you want to use json or jsonb, use jsonb.

SQLite

SQLite supports both bool and bit column types but, unlike pretty much every other database out there, it yields 0 or 1 as the column value instead of false or true. This means that with SQLite alone, you can't just rely on bool or bit columns being treated as truthy/falsey values in Clojure.

You can work around this using a builder that handles reading the column directly as a Boolean:

 $\textbf{com.github.seancorfield/next.jdbc} \quad 1.3.967$

CLJDOC



If you are using plan, you'll most likely be accessing columns by just the label (as a keyword) and avoiding the result set building machinery completely. In such cases, you'll still get bool and bit columns back as 0 or 1 and you'll need to explicitly convert them on a per-column basis since you should know which columns need converting:

See also datafy, nav, and :schema > SQLite for additional caveats on the next.jdbc.datafy namespace when using SQLite.

Friendly SQL Functions

Result Set Builders >

Can you improve this documentation? These fine people already did:

Sean Corfield, Maxim Penzin, Daniel Skarda, Stefan Bleibinhaus, George Peristerakis, Snorre Magnus Davøen & Kimmo Koskinen

Edit on GitHub